

---

# *Computer Vision Classification 1*

WS 2019/2020

Prof. Dr. Simone Frintrop

Computer Vision Group, Department of Informatics  
University of Hamburg, Germany


# Outline

---

- Classification: Problem definition and motivation
- Traditional approach versus deep learning
- Features for Recognition
- Classifier types
- Nearest Neighbor Classifier
- Linear Classification
- Loss functions
- Gradient descent

# Classification

A classifier is an algorithm that assigns to each input image a class label



# Object Categorization

## – Potential Applications



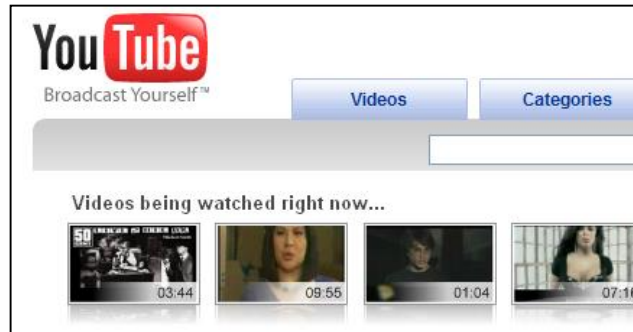
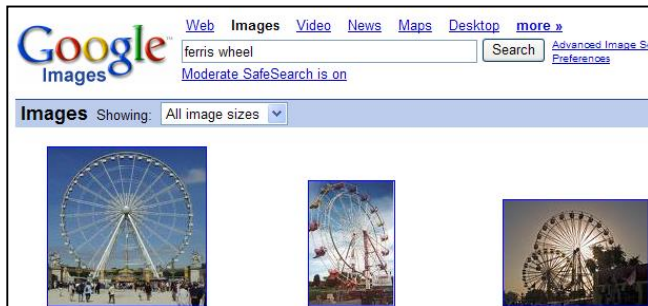
Autonomous robots




Navigation, driver safety



Consumer electronics




Content-based retrieval and analysis for images and videos



Medical image analysis

# How many object categories are there?



1500 – 3000 basic categories (as chair or dog)  
 With subtypes: 10000 to 30000 discriminable  
 objects (as “German shepherd dog”)

[Biederman 1987]



# Identification vs. Categorization

## Identification/Instance Recognition

Find *this particular* object  
(cf. SIFT recognition)



## Categorization/Classification

Find all instances of a  
category/class

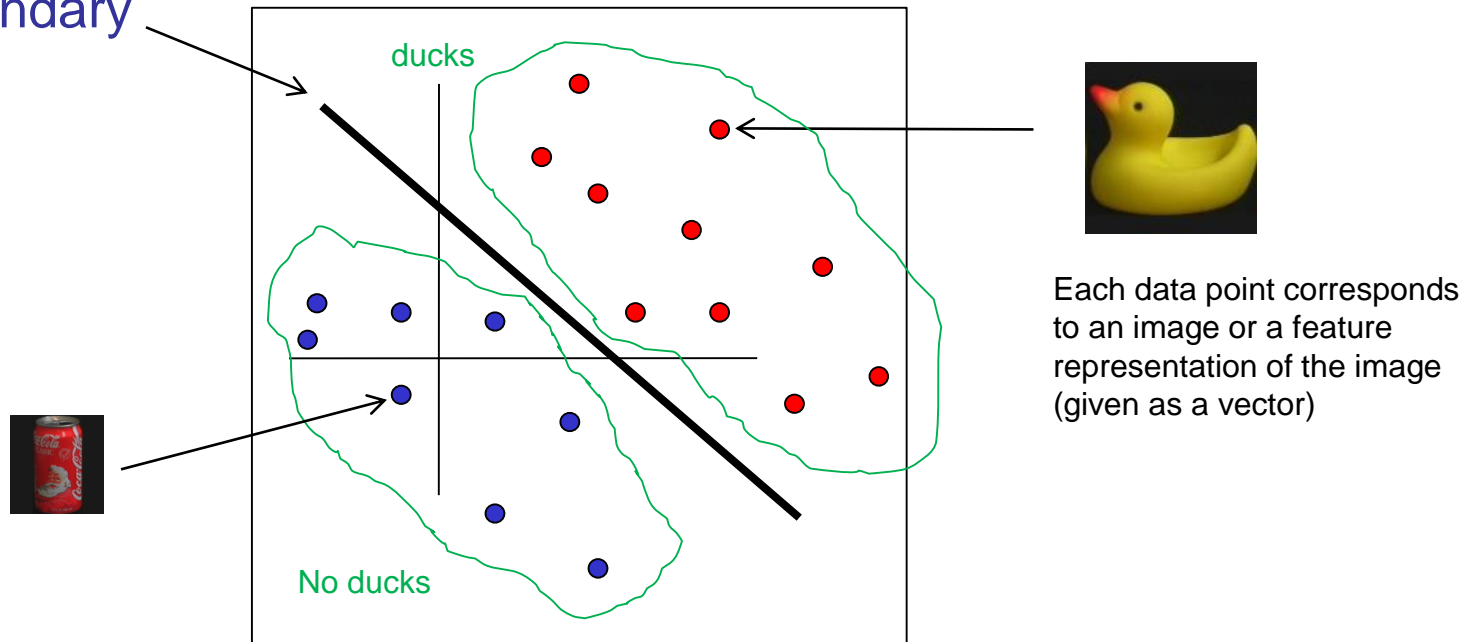


Find any car/cow/...



# Classification

- Data is given as vectors in a feature space
- The classifier assigns **labels** to data points according to a **decision boundary**




- Training a classifier means to determine a decision boundary explicitly or a rule which data point to label with which class (determines the decision boundary implicitly)

# Classification


## ... is easy

if the items within one class are all similar (low *within class variance*) and different from the items of other classes (high *between-class-difference*) in the given feature space



## ... otherwise hard


(if you are lucky, you just chose the wrong feature space)






# Linear vs. Non-Linear Classifiers

**Data separable by linear classifier**




**Data separable only by non-linear classifier**



# *Traditional Recognition Pipeline*

---




# Deep learning

- Learn a *feature hierarchy* all the way from pixels to classifier
- Each layer extracts features from the output of previous layer
- Train all layers jointly



# *Traditional vs. Deep architecture*

## Traditional recognition: “Shallow” architecture




## Deep learning: “Deep” architecture



# *Traditional vs. Deep architecture*

## Traditional recognition: “Shallow” architecture



Lets start with the traditional pipeline



Two problems:

- Which features to select?
- Which classifier to take?

# Features

What object representation do we need?

A vector that has the same length  
for each item/object




Which feature representation we covered is suitable?




# Features

Some feature representations we have seen:

- Color histograms



- HOG features




# Features

---

- Why not SIFT?
- SIFT gives us MANY keypoints per object, and even worse: a different number for each object.  
we do not have a single vector per object.
- But: we can use DenseSIFT or SIFT in a Bag-of-Words approach (see: *Csurka, Dance, Fan, Willamowski & Bray (2004). "Visual categorization with bags of keypoints". Proc. of ECCV International Workshop on Statistical Learning in Computer Vision,*


# Classifiers

## Nearest Neighbor




Berg, Berg, Malik 2005,  
 Chum, Zisserman 2007,  
 Boiman, Shechtman, Irani 2008, ...

## Neural networks




LeCun, Bottou, Bengio, Haffner 1998  
 Rowley, Baluja, Kanade 1998  
 ...

## Boosting




Viola, Jones 2001,  
 Torralba et al. 2004,  
 Opelt et al. 2006,  
 Benenson 2012, ...

## Support Vector Machines



Vapnik, Schölkopf 1995,  
 Papageorgiou, Poggio '01,  
 Dalal, Triggs 2005,  
 Vedaldi, Zisserman 2012

## Randomized Forests



Amit, Geman 1997,  
 Breiman 2001,  
 Lepetit, Fua 2006,  
 Gall, Lempitsky 2009, ...

# *The statistical learning framework*

---

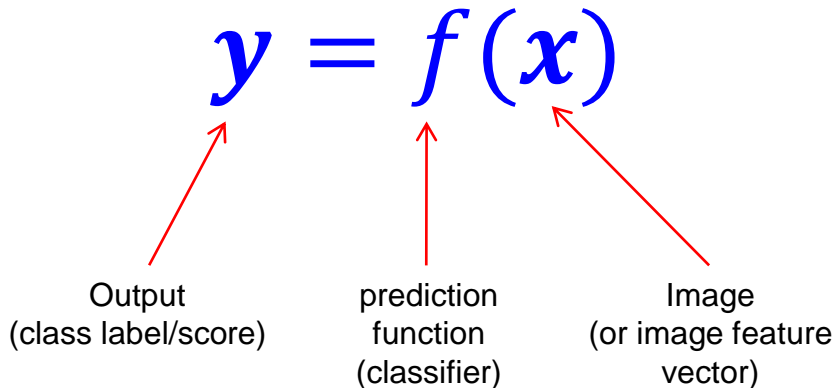
Apply a prediction function to a feature representation of the image to get the desired output:

$$f(\text{apple image}) = \text{"apple"}$$

$$f(\text{tomato image}) = \text{"tomato"}$$

$$f(\text{cow image}) = \text{"cow"}$$

# The statistical learning framework



$f(\text{apple image})$	= "apple"
$f(\text{tomato image})$	= "tomato"
$f(\text{cow image})$	= "cow"

- **Training:** given a *training set* of labeled examples  $\{(x_1, y_1), \dots, (x_N, y_N)\}$ , estimate the prediction function  $f$  by minimizing the prediction error on the training set
- **Testing:** apply  $f$  to an unseen *test example*  $x$  and output the predicted value  $y = f(x)$

# *A simple linear classifier*

A simple linear function (1D input):

$$y = wx + b$$


It separates two classes (binary classification)

Multi-dimensional input  $\mathbf{x}$  (e.g. image with  $n = x*y$  pixels)

$$y = \mathbf{w}^T \mathbf{x} + b$$

Note:  $\mathbf{w}$ ,  $\mathbf{x}$  are vectors now

This is the same as:  $y = \sum_{i=1}^n w_i x_i + b$





# Parametric approach


Before:  $y = f(x) = \mathbf{w}^T \mathbf{x} + b$

Consider now the parameters not as part of the function, but as input that can be trained (parametric approach):

$$y = f(\mathbf{x}, \mathbf{w}, b)$$

Find function  $f$  with a set of parameters  $\mathbf{w}$  and  $b$  that maps images to classes

We want to learn the  $\mathbf{w}$  and  $b$  that optimize this function.



# *The statistical learning framework*

---

Find function  $f$  with a set of parameters  $w$  and  $b$  that maps images to classes:

$$f(\text{apple image}, W, b) = \text{"apple"}$$

$$f(\text{tomato image}, W, b) = \text{"tomato"}$$

$$f(\text{cow image}, W, b) = \text{"cow"}$$

# Classifiers

The following questions remain:

- What function can we choose for  $f$ ?  
(Which classifier shall we use?)
- How can we measure how good the classifier is?  
(Define a loss function)
- How can we optimize the classifier?  
(Find better  $W, b$ )

$$f(\text{apple}, W, b) = \text{"apple"}$$

$$f(\text{tomato}, W, b) = \text{"tomato"}$$

$$f(\text{cow}, W, b) = \text{"cow"}$$

# Classifiers

The following questions remain:

- What function can we choose for  $f$ ?  
(Which classifier shall we use?)
- How can we measure how good the classifier is?  
(Define a loss function)
- How can we optimize the classifier?  
(Find better  $W, b$ )

$$f(\text{apple}, W, b) = \text{"apple"}$$

$$f(\text{tomato}, W, b) = \text{"tomato"}$$

$$f(\text{cow}, W, b) = \text{"cow"}$$

# Classifiers

- What function can we choose for  $f$ ?
- Well-known classifiers:
  - Logistic regression
  - Support Vector Machine
  - Boosting
  - Random Forests
  - (Deep) Neural Networks

$$\begin{aligned} f(\text{apple}, W, b) &= \text{"apple"} \\ f(\text{tomato}, W, b) &= \text{"tomato"} \\ f(\text{cow}, W, b) &= \text{"cow"} \end{aligned}$$

We will cover here mainly Deep Neural Networks,  
but let's continue now with the big picture

# Classifiers

- Given a function  $f$  that maps images to classes, and given a set of parameters  $W$  and  $b$ :  
What output do we get?
- Usually a **classification score** that measures how certain the classifier is for each class.  
The score is a vector. Size: number of classes
- Binary case:

$$f(\text{tomato}, w, b): \begin{pmatrix} 0.88 \\ 0.12 \end{pmatrix} \begin{matrix} \text{tomato} \\ \text{no tomato} \end{matrix}$$



# Classifiers

- Given a function  $f$  that maps images to classes, and given a set of parameters  $W$  and  $b$ :  
What output do we get?
- Usually a **classification score** that measures how certain the classifier is for each class.  
The score is a vector. Size: number of classes
- Multi-class case:

$$f(\text{🍅}, W, b): \begin{pmatrix} 0.54 \\ 4.78 \\ -2.56 \\ \vdots \\ \vdots \\ \vdots \end{pmatrix} \begin{matrix} \text{Car} \\ \text{Tomato} \\ \text{Chair} \end{matrix}$$

(whether the values sum up to 1 or not depends on the classifier)

# A simple linear classifier

A simple linear function (1D input):

$$y = wx + b$$


It separates two classes (binary classification)

Multi-dimensional input  $\mathbf{x}$  (e.g. image with  $n = x*y$  pixels)

$$y = \mathbf{w}^T \mathbf{x} + b$$

Note:  $\mathbf{w}$ ,  $\mathbf{x}$  are vectors now

This is the same as:  $y = \sum_{i=1}^n w_i x_i + b$



# Multi-Class Linear Classifier

Let's extend our linear classifier to multi-class:

Up to now we had:

$$y = \mathbf{w}^T \mathbf{x} + b$$

(multi-dimensional input  $\mathbf{x} \Rightarrow$  vectors  $\mathbf{x}$  and  $\mathbf{w}$ )

Now we need: multi-dimensional output  $\mathbf{y}$ :

$$\mathbf{y} = f(\mathbf{x}, \mathbf{W}, \mathbf{b}) = \mathbf{W}\mathbf{x} + \mathbf{b}$$

where  $\mathbf{x}$ ,  $\mathbf{b}$  and  $\mathbf{y}$  are vectors and  $\mathbf{W}$  is a matrix

$$\begin{pmatrix} 0.54 \\ 4.78 \\ -2.56 \\ . \\ . \\ . \end{pmatrix} \begin{matrix} \text{Car} \\ \text{Tomato} \\ \text{Chair} \end{matrix}$$

# Example

---

Image from  
CIFAR-10  
dataset:



**[32x32x3]**

= 3072 pixels

What are the dimensions of the vectors/matrix?

$$\mathbf{y} = f(\mathbf{x}, \mathbf{W}, \mathbf{b}) = \mathbf{W}\mathbf{x} + \mathbf{b}$$

CIFAR-10 dataset: 60000 color images of 32x32 pixels in 10 classes

# Example



**[32x32x3]**

array of numbers 0...1


$$\underbrace{f(x, W)}_{10 \times 1} = \underbrace{W}_{10 \times 3072} \underbrace{x}_{3072 \times 1} + \underbrace{(+b)}_{10 \times 1}$$

**10** numbers,  
indicating class  
scores

parameters, or “weights”

# Linear classifier


Example with an image with 4 pixels, and 3 classes (cat/dog/ship)



Each row of  $W$  is a linear classifier.



# Interpreting a Linear Classifier



$$f(x_i, W, b) = Wx_i + b$$




**[32x32x3]**  
array of numbers 0...1  
(3072 numbers total)

Each row of **W** is a classifier for one of the classes, visualized here as line

# *Multi-Class Classification*

Example of the decision boundaries learnable with a linear classifier for multiple classes:



(in higher dimensions, the lines are hyperplanes)

<https://shapeofdata.wordpress.com/2013/06/04/multi-class-classification/>

# Multi-Class Linear Classifier

---

Up to now: equations for *one* training image  $\mathbf{x}$ :

$$\mathbf{y} = f(\mathbf{x}, \mathbf{W}, \mathbf{b}) = \mathbf{W}\mathbf{x} + \mathbf{b}$$

where  $\mathbf{x}$ ,  $\mathbf{b}$  and  $\mathbf{y}$  are vectors and  $\mathbf{W}$  is a matrix

Extension to multiple training images:

$$\mathbf{Y} = f(\mathbf{X}, \mathbf{W}, \mathbf{B}) = \mathbf{W}\mathbf{X} + \mathbf{B}$$

where  $\mathbf{Y}$ ,  $\mathbf{X}$ ,  $\mathbf{W}$ , and  $\mathbf{B}$  are all matrices

=> matrix-oriented programming languages like Matlab or Python are very convenient to compute this

# Equation overview

$$y = wx + b$$

$$y = \mathbf{w}^T \mathbf{x} + b$$

$$\mathbf{y} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{Y} = \mathbf{W}\mathbf{X} + \mathbf{B}$$

Classification Type	Input x	Parameters w/W	Bias b	Output y
Binary (two classes)	scalar	scalar	scalar	scalar
Binary (two classes)	vector (e.g. image)	vector	scalar	scalar
Multi-class	vector (e.g. image)	matrix	vector	vector
Multi-class	matrix (several images)	matrix	matrix	matrix

# Classifiers

The following questions remain:

- What function can we choose for  $f$ ?  
(Which classifier shall we use?)
- How can we measure how good the classifier is?  
(Define a loss function)
- How can we optimize the classifier?  
(Find better  $W, b$ )

$$f(\text{apple}, W, b) = \text{"apple"}$$

$$f(\text{tomato}, W, b) = \text{"tomato"}$$

$$f(\text{cow}, W, b) = \text{"cow"}$$

# How good is our function?

- Given a function  $f$  that maps images to classes, and given a set of parameters  $W$  and  $b$ .

$$y = f(\mathbf{x}, \mathbf{W}, \mathbf{b})$$

$$\begin{aligned} f(\text{apple}, W, b) &= \text{"apple"} \\ f(\text{tomato}, W, b) &= \text{"tomato"} \\ f(\text{cow}, W, b) &= \text{"cow"} \end{aligned}$$

- How good is  $f$ ?
- We can compute the **loss** over the training data

# Loss functions

- A **loss function** (also cost function) measures how good your result is
- It compares the given output with the target output

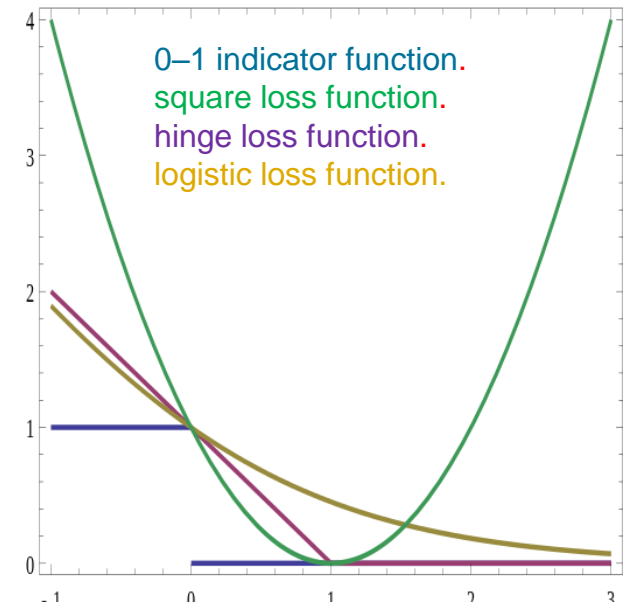
			Loss	
$f(\text{cow}, W, b)$	= "cow"	correct 😊	0	Don't punish the classifier for this
$f(\text{apple}, W, b)$	= "tomato"	wrong 😞	1	Punish the classifier for this
$f(\text{orange}, W, b)$	= "cow"	wrong 😞	1	Punish the classifier for this

Average loss: 2/3

- This is the **0-1 loss**. It is good for illustration purposes, but usually not useful (non-convex, non smooth)

# Loss functions

- A **loss function** (also cost function) measures how good your result is
- It compares the given output with the target output
- Popular loss functions:
  - 0-1 loss/0-1 indicator function
  - Mean square loss
  - SVM loss (hinge loss)
  - Logistic loss
  - Cross-entropy loss
- The choice of your loss function is not arbitrary. It has to fit to the model and the problem!



[Wikipedia: loss functions for classification]



# Full Loss

---

Given a loss function that computes the loss  $L_i$  for a given input image, we obtain the full loss  $L$  by averaging over all  $N$  training images:

$$L = \frac{1}{N} \sum_{i=1}^N L_i + R(W)$$

where  $R(W)$  is a regularization term that prevents overfitting (regularization will be covered in more detail in summer term lecture „Computer Vision 2“)

# Classifiers

The following questions remain:

- What function can we choose for  $f$ ?  
(Which classifier shall we use?)
- How can we measure how good the classifier is?  
(Define a loss function)
- **How can we optimize the classifier? (Find better  $W, b$ )**

$$f(\text{apple}, W, b) = \text{"apple"}$$

$$f(\text{tomato}, W, b) = \text{"tomato"}$$

$$f(\text{cow}, W, b) = \text{"cow"}$$

# Classifiers

How can we optimize the classifier?

$$\begin{aligned} f(\text{apple}, W, b) &= \text{"apple"} \\ f(\text{tomato}, W, b) &= \text{"tomato"} \\ f(\text{cow}, W, b) &= \text{"cow"} \end{aligned}$$

Q: what is a „better“ classifier?


A: it has a lower loss => minimize the loss

Q: What can we change to obtain a lower loss?

A: the parameters  $W$  and  $b$

# Gradient descent

- Plot loss function dependent on the weights:
- A convex loss function  $L(w)$  for a single weight  $w$ :



Find the minimum by  
walking downwards:  
*Gradient descent*

We follow the slope!  
The slope is given by  
the derivative of  $L(w)$ :

$$\frac{dL(w)}{dw}$$

# Gradient descent

$L(w)$



## Gradient descent:

- Initialize  $w$  randomly.
- Repeat until convergence


$$\left\{ \begin{array}{l} w := w - \alpha \frac{dL(w)}{dw} \end{array} \right\}$$

Learning rate

This means: update weight value  $w$  by subtracting a portion of the derivative

# Gradient descent

If we have two weights with loss function  $L(w_1, w_2)$ :



Initialize  $w$  randomly

Repeat until convergence:


$$\left\{ \begin{array}{l} w_1 := w_1 - \alpha \frac{\partial L(w_1, w_2)}{\partial w_1} \\ w_2 := w_2 - \alpha \frac{\partial L(w_1, w_2)}{\partial w_2} \end{array} \right\}$$

**Note:** the result of updating both weights means to step into the direction of the negative gradient!

# *Optimizing the loss*

For a non-linear function (like a multi-layer neural network) the loss function is usually more complicated:


For 2 parameters (weights):



And we have usually thousands/millions of weights (dimensions)!

# Optimizing the loss

- But we can still use gradient descent



- Note: does only converge to *local* optimum
- But: it often works very well in practice (local optima often are already a good solution)
- Gradient descent is the standard optimization in deep learning



# Brief summary

- Choose a data representation  $x$  (features or simply pixels)
- Select a classifier that computes  $f(x, W, b)$  (the score)
- Compute the loss
- Optimize  $f$  by finding better parameters  $W, b$  that minimize the loss

$$f(\text{apple}, W, b) = \text{"apple"}$$

$$f(\text{tomato}, W, b) = \text{"tomato"}$$

$$f(\text{cow}, W, b) = \text{"cow"}$$

- Next slide set: deep learning to compute  $f$

# Primary literature

---

You find most of these contents in most machine learning books, however not always related to images and seldomly in condensed form. Read e.g. the related parts from:

- Gonzales/Woods 2018: relevant parts from chapter 13
- Goodfellow et al, “Deep Learning”, 2016: relevant parts from chapter 5 (machine learning basics)
- Course notes: [CS231n Convolutional Neural Networks for Visual Recognition](#) (lectures 1-3), from Stanford lecture of Fei Fei Li, Andrej Karpathy and Justin Johnson  
<http://cs231n.github.io/>  
(Corresponding video lectures on Youtube also recommended!)

# *Secondary literature*

---

- Video lecture: “Deep Learning by Andrew Ng”  
<https://www.youtube.com/playlist?list=PLBAGcD3siRDguYyYYzhVwZ3tLvOyyG5k6K>
- Biederman, I. (1987). Recognition-by-components: a theory of human image understanding. *Psychological review*, 94(2):115.
- Russakovsky, O., Deng, J., Su, H., Krause, J., Satheesh, S., Ma, S., Huang, Z., Karpathy, A., Khosla, A., Bernstein, M., et al. (2015). ImageNet large scale visual recognition challenge. *International Journal of Computer Vision*, 115(3):211–252.
- Boiman/Shechtman/Irani: “In defense of Nearest-Neighbor Based Image Classification”, CVPR 2008